

Delhi temperature timeseries data SARIMA model fitting

Here we are doing a time-series analysis of the daily temperature in the city of Delhi for the period of Jan 01, 1995 to Dec 31, 2019.

Reading the data

```
DelhiTemp <- read.csv(file = '../Data/Delhi_temperature_1995_2020.csv')
```

Viewing the data

```
head(DelhiTemp)
```


```
##   Month Day Year Temperature
## 1     1   1 1995         50.7
## 2     1   2 1995         52.1
## 3     1   3 1995         53.8
## 4     1   4 1995         53.7
## 5     1   5 1995         54.5
## 6     1   6 1995         54.3
```

Creating a date column from Month, Day and Year columns.

```
DelhiTemp$Date <- as.Date(with(DelhiTemp, paste(Year, Month, Day, sep="-")), "%Y-%m-%d")
```

Temperature vs. Date plot

```
Temp <- DelhiTemp$Temperature
Date <- DelhiTemp$Date
plot(Temp~Date)
```



Our

data shows clear yearly seasonality and there are some very large outliers (temperature ~ 100) possibly because of missing data for those dates.

Outlier detection and missing value treatment

Identifying outliers and replacing them with backfilled values:

```
DelhiTemp$Temperature[DelhiTemp$Temperature < 0] <- NA
which(is.na(DelhiTemp$Temperature))
```


```
## [1] 88 1324 1370 1454 1455 1460 1461 1471 1580 1581 2310 2726 2727 2728 2729
## [16] 2856 4448 4449 4523 4623 4678 4724 5016 5175 5213 5586 5587 5588 5590 6004
## [31] 6005 6006 6007 6008 6710 6977 7670 7740 7777 7796 7935 7986 8145 8146 8405
## [46] 8661 8668 8719 8720 8722 8789 8880 8881 8883 8886 8902 8903 8904 8905
```

```
DelhiTemp$Temperature <- na.locf(DelhiTemp$Temperature, fromLast = TRUE)
which(is.na(DelhiTemp$Temperature))
```

```
## integer(0)
```

Updated Temperature vs. Date plot

```
Temp <- DelhiTemp$Temperature
Date <- DelhiTemp$Date
plot(Temp~Date)
```



Binning

Binning temperatures to their mean monthly values

```
DelhiMonthlyTemp <- DelhiTemp %>% group_by(Year,Month) %>% summarize(Temperature = mean(Temperature))
```

```
## `summarise()` has grouped output by 'Year'. You can override using the `.groups` argument.
```

```
head(DelhiMonthlyTemp);
```


```
## # A tibble: 6 x 3
## # Groups:   Year [1]
##   Year Month Temperature
##   <int> <int>      <dbl>
## 1  1995     1        55.5
## 2  1995     2        62.4
## 3  1995     3        70.5
## 4  1995     4        83.1
## 5  1995     5        93.2
## 6  1995     6        95.8
```

Setting a new date column with just the month and year info

```
DelhiMonthlyTemp$Date <- as.yearmon(paste(DelhiMonthlyTemp$Year, DelhiMonthlyTemp$Month), "%Y %m")
```

Monthly mean temperature vs. Date (Month) plot

```
Temp <- DelhiMonthlyTemp$Temperature
Date <- DelhiMonthlyTemp$Date
plot(Temp~Date, type="l")
```



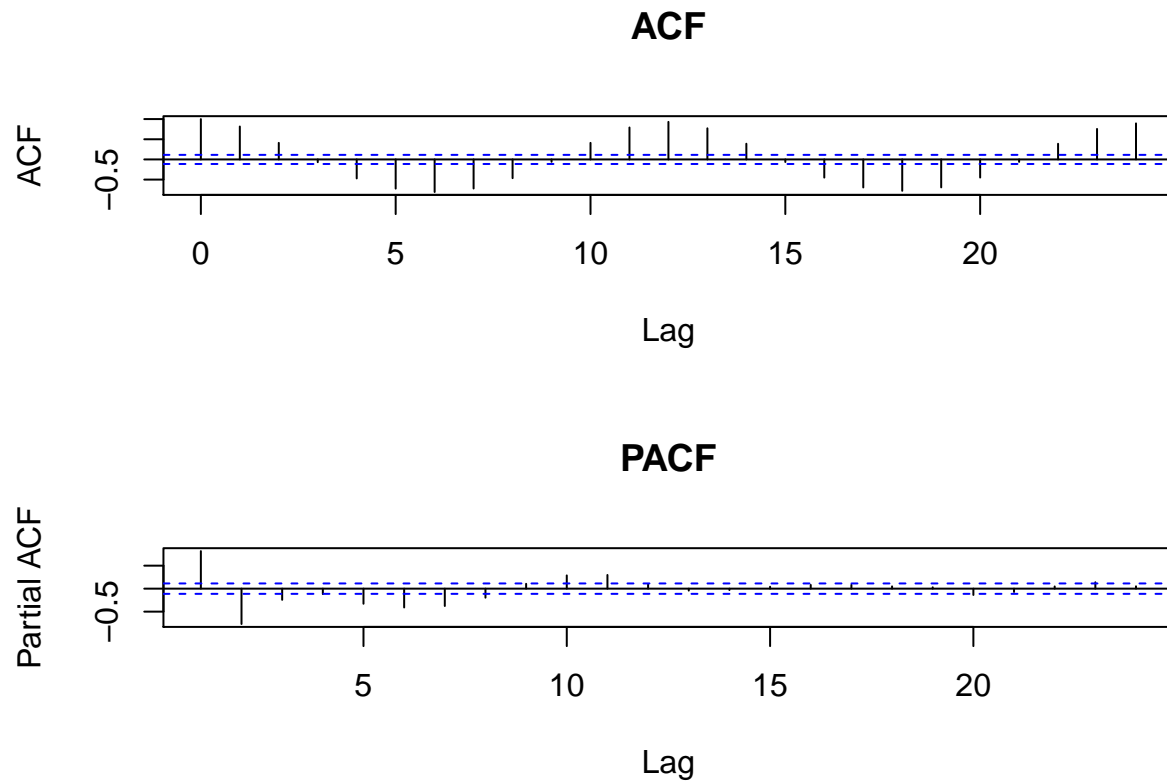
Once again we see that the data shows clear seasonality but no variation in variance so differencing should be enough for taking care of the trend.

ACF and PACF plots

Let's first look at the auto-correlation function and the partial auto-correlation function plots.

```
Temp <- DelhiMonthlyTemp$Temperature
par(mfrow=c(2,1))
```

```
acf(Temp,main='ACF')
pacf(Temp,main='PACF');
```




We see that the ACF also shows clear seasonality.

Guessing the right orders for (S)ARIMA model fitting

1. Differencing orders (d, D)

Non-seasonal and seasonal differencing -> `diff(diff(data),12)`

```
plot(diff(diff(Temp),12),type="l")
```



the plot shows almost no-trends except for a few large peaks at the center which may be outliers => $d=1$, $D=1$. We can also use the AD Fuller test for checking the stationarity

```
#diff_data <- diff(diff(Temp),12)
#adf.test(diff_data)
```


The test confirms that the data is stationary.

2. orders for the auro-regressive (AR) and Moving Average (MA) terms i.e. p and q


ACF and PACF for differenced data

```
par(mfrow=c(2,1))
acf(diff(diff(Temp),12),main='differnced data ACF',50)
pacf(diff(diff(Temp),12),main='differnced data PACF',50);
```

differrnced data ACF



differrnced data PACF



The ACF plot shows significant correlations at lag=1,11 and 12 while the PACF shows significant correlation for lag=1,11 and 12. The correlations at later parts may be due to seasonality.

Finding best parameters using

1. Grid Search

Trying for different values of p,q,P,Q and note down AIC, SSE and p-value (for Ljung-Box-test). We want high p-values and small AIC and SSE using parsimony principle (simpler the better) while searching.

```
d=1; DD=1; per=12
```

```
for(p in 1:2){
  for(q in 1:2){
    for(i in 1:6){
      for(j in 1:3){
        if(p+d+q+i+DD+j<=10){

          model<-arima(x=Temp, order = c((p-1),d,(q-1)), seasonal = list(order=c((i-1),DD,(j-1)), period=per))

          pval<-Box.test(model$residuals, lag=log(length(model$residuals)))

          sse<-sum(model$residuals^2)

          cat(p-1,d,q-1,i-1,DD,j-1,per, 'AIC=', model$aic, ' SSE=', sse, ' p-VALUE=', pval$p.value, '\n')

        }
      }
    }
  }
}
```

```
}
}
```

```
## 0 1 0 0 1 0 12 AIC= 1499.338 SSE= 3099.001 p-VALUE= 8.952575e-08
## 0 1 0 0 1 1 12 AIC= 1320.604 SSE= 1476.127 p-VALUE= 8.894564e-07
## 0 1 0 0 1 2 12 AIC= 1320.109 SSE= 1502.587 p-VALUE= 1.212609e-06
## 0 1 0 1 1 0 12 AIC= 1392.531 SSE= 2087.502 p-VALUE= 1.215459e-06
## 0 1 0 1 1 1 12 AIC= 1319.835 SSE= 1500.886 p-VALUE= 1.226718e-06
## 0 1 0 1 1 2 12 AIC= 1320.854 SSE= 1488.524 p-VALUE= 1.356503e-06
## 0 1 0 2 1 0 12 AIC= 1370.082 SSE= 1901.339 p-VALUE= 7.147622e-07
## 0 1 0 2 1 1 12 AIC= 1321.25 SSE= 1485.317 p-VALUE= 1.148136e-06
## 0 1 0 2 1 2 12 AIC= 1322.845 SSE= 1486.992 p-VALUE= 1.336722e-06
## 0 1 0 3 1 0 12 AIC= 1349.824 SSE= 1740.855 p-VALUE= 4.246554e-06
## 0 1 0 3 1 1 12 AIC= 1322.842 SSE= 1494.211 p-VALUE= 1.265004e-06
## 0 1 0 4 1 0 12 AIC= 1344.303 SSE= 1686.293 p-VALUE= 3.323896e-06
## 0 1 1 0 1 0 12 AIC= 1409.009 SSE= 2237.481 p-VALUE= 0.02513918
## 0 1 1 0 1 1 12 AIC= 1238.17 SSE= 1070.666 p-VALUE= 0.01793177
## 0 1 1 0 1 2 12 AIC= 1239.052 SSE= 1060.471 p-VALUE= 0.01667548
## 0 1 1 1 1 0 12 AIC= 1319.848 SSE= 1607.407 p-VALUE= 0.01328627
## 0 1 1 1 1 1 12 AIC= 1239.027 SSE= 1060.696 p-VALUE= 0.01664074
## 0 1 1 1 1 2 12 AIC= 1240.96 SSE= 1061.698 p-VALUE= 0.01650034
## 0 1 1 2 1 0 12 AIC= 1295.577 SSE= 1453.315 p-VALUE= 0.007799349
## 0 1 1 2 1 1 12 AIC= 1240.994 SSE= 1061.651 p-VALUE= 0.01664093
## 0 1 1 3 1 0 12 AIC= 1282.807 SSE= 1369.994 p-VALUE= 0.001596375
## 1 1 0 0 1 0 12 AIC= 1460.192 SSE= 2683.762 p-VALUE= 5.570336e-05
## 1 1 0 0 1 1 12 AIC= 1288.522 SSE= 1281.156 p-VALUE= 2.515287e-05
## 1 1 0 0 1 2 12 AIC= 1289.366 SSE= 1269.078 p-VALUE= 2.049262e-05
## 1 1 0 1 1 0 12 AIC= 1361.389 SSE= 1860.666 p-VALUE= 7.996062e-05
## 1 1 0 1 1 1 12 AIC= 1289.222 SSE= 1272.065 p-VALUE= 2.021527e-05
## 1 1 0 1 1 2 12 AIC= 1290.516 SSE= 1267.385 p-VALUE= 2.03264e-05
## 1 1 0 2 1 0 12 AIC= 1338.392 SSE= 1691.143 p-VALUE= 9.26347e-05
## 1 1 0 2 1 1 12 AIC= 1290.295 SSE= 1270.33 p-VALUE= 2.401602e-05
## 1 1 0 3 1 0 12 AIC= 1324.173 SSE= 1585.588 p-VALUE= 6.74795e-05
## 1 1 1 0 1 0 12 AIC= 1385.017 SSE= 2016.254 p-VALUE= 0.9964005
## 1 1 1 0 1 1 12 AIC= 1221.998 SSE= 978.4572 p-VALUE= 0.9779703
## 1 1 1 0 1 2 12 AIC= 1223.462 SSE= 972.9903 p-VALUE= 0.9678169
## 1 1 1 1 1 0 12 AIC= 1298.865 SSE= 1459.958 p-VALUE= 0.9675623
## 1 1 1 1 1 1 12 AIC= 1223.441 SSE= 972.6935 p-VALUE= 0.9671796
## 1 1 1 2 1 0 12 AIC= 1277.296 SSE= 1331.71 p-VALUE= 0.9371032
```

2. Using auto.arima()

```
y <- msts(Temp, seasonal.periods=c(12))
auto.arima(y, d = 1, D = 1, max.p = 5, max.q = 5, max.P = 5, max.Q = 5, max.order = 10, start.p =

## Series: y
## ARIMA(5,1,0)(5,1,0)[12]
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ar5      sar1      sar2      sar3
##      -0.5142  -0.4444  -0.3373  -0.2261  -0.1488  -0.8628  -0.6184  -0.4755
## s.e.   0.0606   0.0653   0.0695   0.0676   0.0603   0.0625   0.0823   0.0867
##          sar4      sar5
##      -0.2962  -0.1132
## s.e.   0.0850   0.0654
```

```
##
## sigma^2 estimated as 4.656:  log likelihood=-628.94
## AIC=1279.88   AICc=1280.84   BIC=1320.13
```


Best-model

For some reason auto-arima is unable to reproduce the minimum value of AIC which was found in the grid-search method. From the grid-search the lowest AIC of 1221.998 is found for a 1,1,1,0,1,1,12 SARIMA model which also has a large enough p-value.

SARIMA(1,1,1,0,1,1,12) fitting results

```
sarima(Temp,1,1,1,0,1,1,12);
```

```
## initial  value 1.191317
## iter    2 value 0.898750
## iter    3 value 0.873868
## iter    4 value 0.851006
## iter    5 value 0.826823
## iter    6 value 0.797812
## iter    7 value 0.748496
## iter    8 value 0.742428
## iter    9 value 0.739291
## iter   10 value 0.737991
## iter   11 value 0.735174
## iter   12 value 0.733637
## iter   13 value 0.733464
## iter   14 value 0.733395
## iter   15 value 0.733379
## iter   16 value 0.733377
## iter   17 value 0.733375
## iter   18 value 0.733375
## iter   18 value 0.733375
## final   value 0.733375
## converged
## initial  value 0.716771
## iter    2 value 0.699040
## iter    3 value 0.696613
## iter    4 value 0.696121
## iter    5 value 0.696043
## iter    6 value 0.696041
## iter    7 value 0.696040
## iter    7 value 0.696040
## iter    7 value 0.696040
## final   value 0.696040
## converged
```

```
## $fit
##
## Call:
## stats::arima(x = xdata, order = c(p, d, q), seasonal = list(order = c(P, D,
##     Q), period = S), include.mean = !no.constant, transform.pars = trans, fixed = fixed,
##     optim.control = list(trace = trc, REPORT = 1, reltol = tol))
##
## Coefficients:
##      ar1      ma1      sma1
##    0.3468 -1.0000 -0.9999
## s.e.  0.0563  0.0386  0.0874
##
## sigma^2 estimated as 3.409:  log likelihood = -607,  aic = 1222
##
## $degrees_of_freedom
## [1] 284
##
## $ttable
##      Estimate      SE  t.value p.value
## ar1    0.3468 0.0563   6.1606      0
## ma1   -1.0000 0.0386  -25.8977      0
## sma1  -0.9999 0.0874  -11.4404      0
##
## $AIC
## [1] 4.100663
##
```

```
## $AICc
## [1] 4.100937
##
## $BIC
## [1] 4.149784
```

We see that the residuals are almost stationary and there is not much correlation left in the lags of the residuals, also the p-value of the Ljung-Box test is significant at almost all lags.

Forecasting using the best-model

```
model<-arima(x=Temp, order = c(1,1,1), seasonal = list(order=c(0,1,1), period=per))
par(mfrow=c(1,1))
plot(forecast(model,12)); # forecasting for the 12 months after the end of the dataset
```

Forecasts from ARIMA(1,1,1)(0,1,1)[12]

